



# OPERATING SYSTEM

## Threads & Concurrency

---

**Dr Rahul Nagpal**  
Computer Science

# OPERATING SYSTEM

---

## Threads & Concurrency

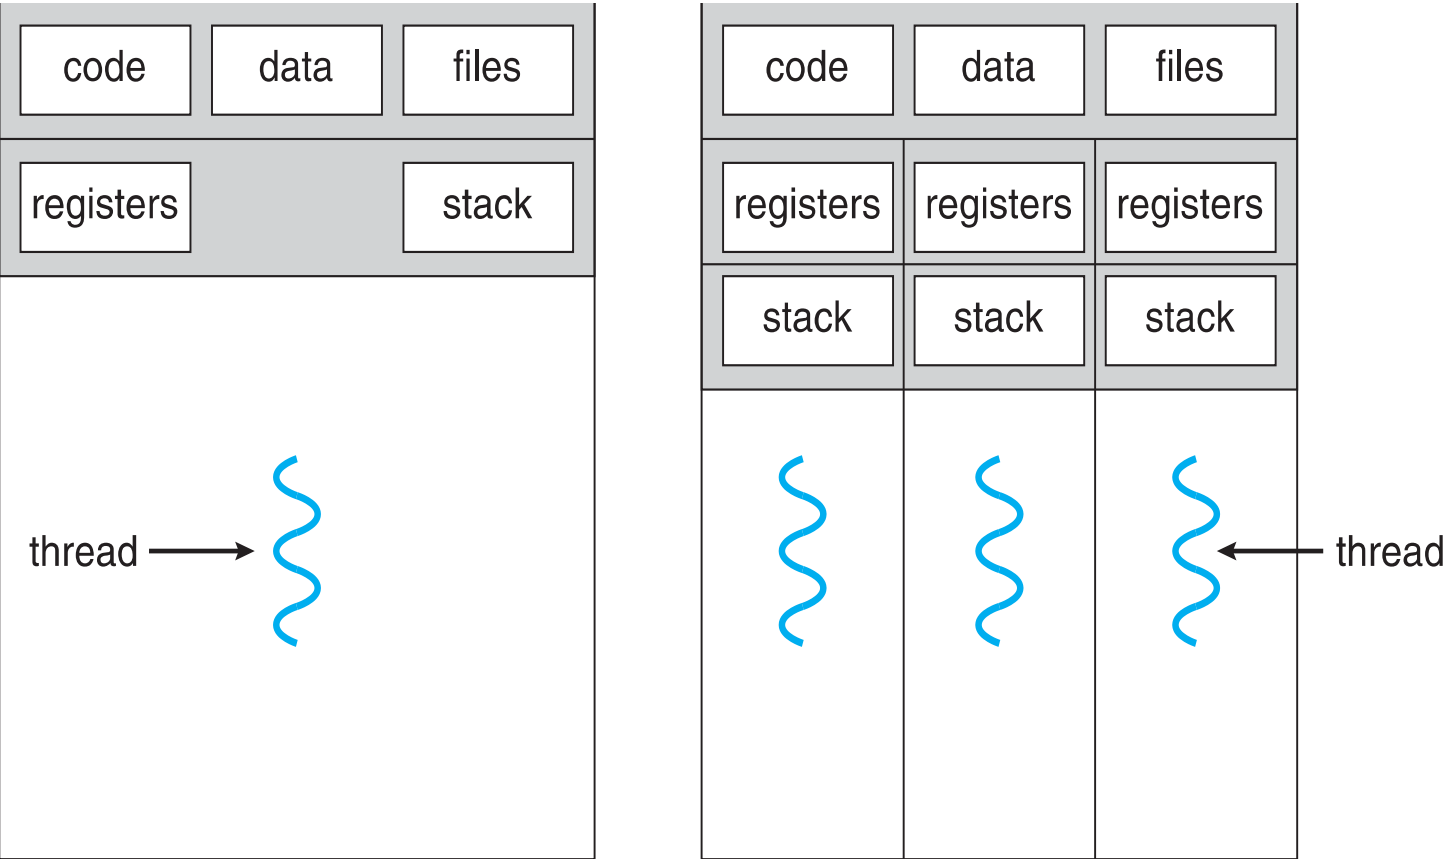
**Dr. Rahul Nagpal**  
Computer Science

- Threads and Concurrency
- Multicore Programming
- Multithreading Models
- Thread Scheduling

- Process creation is heavy-weight while thread creation is light-weight
- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks in application can be implemented by threads
  - Update display
  - Fetch data
  - Spell checking
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

# OPERATING SYSTEMS

## Single and Multithreaded Processes

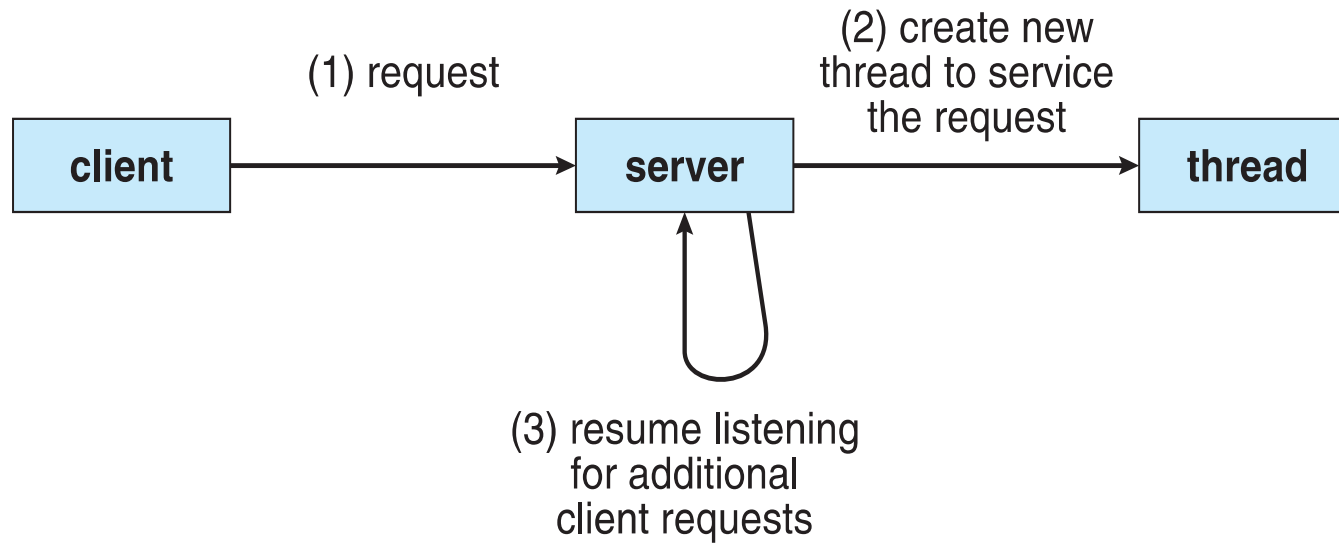


single-threaded process

multithreaded process

# OPERATING SYSTEMS

## Multithreaded Server Architecture



**Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces

**Resource Sharing** – threads share resources of process, easier than shared memory or message passing

**Economy** – cheaper than process creation, thread switching lower overhead than context switching

**Scalability** – process can take advantage of multiprocessor architectures

**Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:

- Dividing activities**

- Balance**

- Data splitting**

- Data dependency**

- Testing and debugging**

***Parallelism*** implies a system can perform more than one task simultaneously

***Concurrency*** supports more than one task making progress  
Single processor / core, scheduler providing concurrency



### Types of parallelism

**Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each

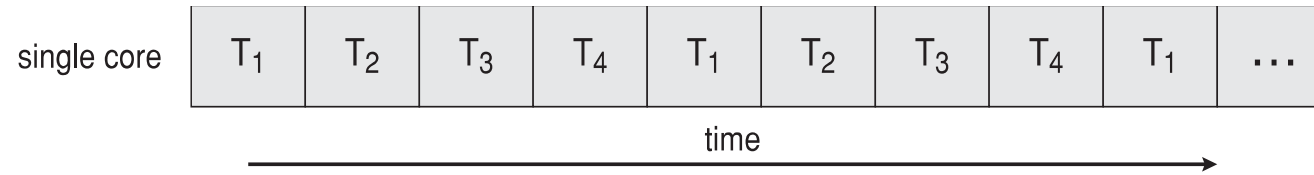
**Task parallelism** – distributing threads across cores, each thread performing unique operation

As # of threads grows, so does architectural support for threading

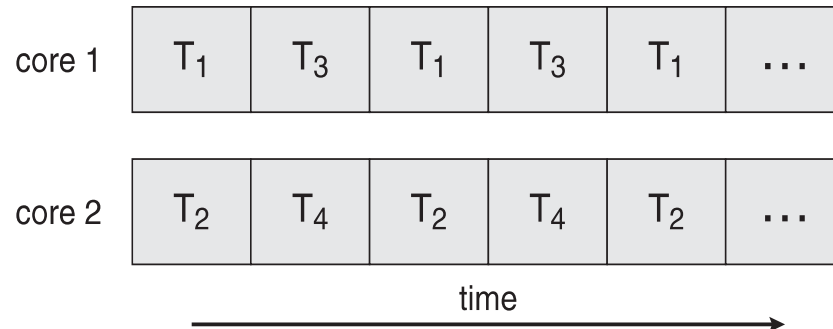
CPUs have cores as well as *hardware threads*

Consider Oracle SPARC T4 with 8 cores, and 8 hardware threads per core

- **Concurrent execution on single-core system:**



- **Parallelism on a multi-core system:**



- **User threads** - management done by user-level threads library
- Three primary thread libraries:
  - POSIX **Pthreads**
  - Windows threads
  - Java threads
- **Kernel threads** - Supported by the Kernel
- Examples – virtually all general purpose operating systems, including:
  - Windows
  - Solaris
  - Linux
  - Tru64 UNIX
  - Mac OS X

# OPERATING SYSTEMS

## User Threads and Kernel Threads

---

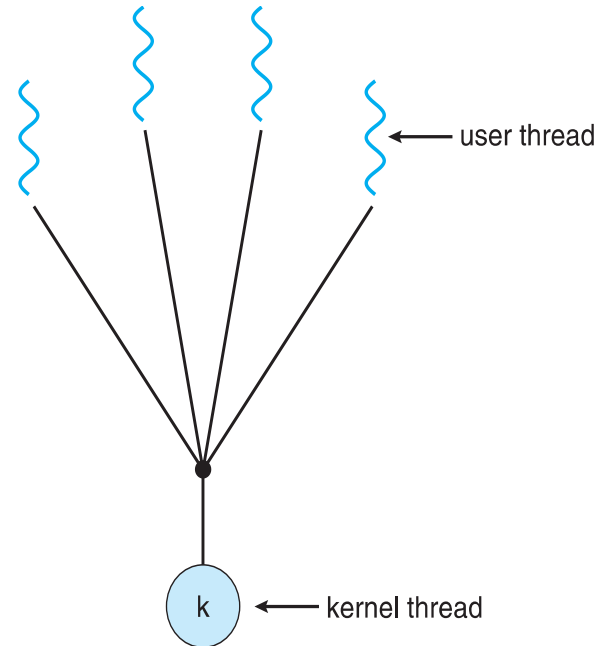
- Many-to-One
- One-to-One
- Many-to-Many



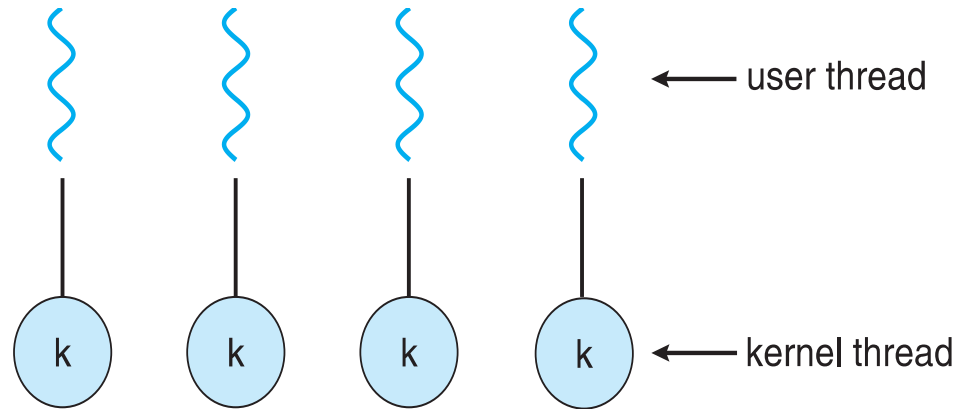
# OPERATING SYSTEMS

## Many-to-One

- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on muticore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
  - **Solaris Green Threads**
  - **GNU Portable Threads**



- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
  - Windows
  - Linux
  - Solaris 9 and later

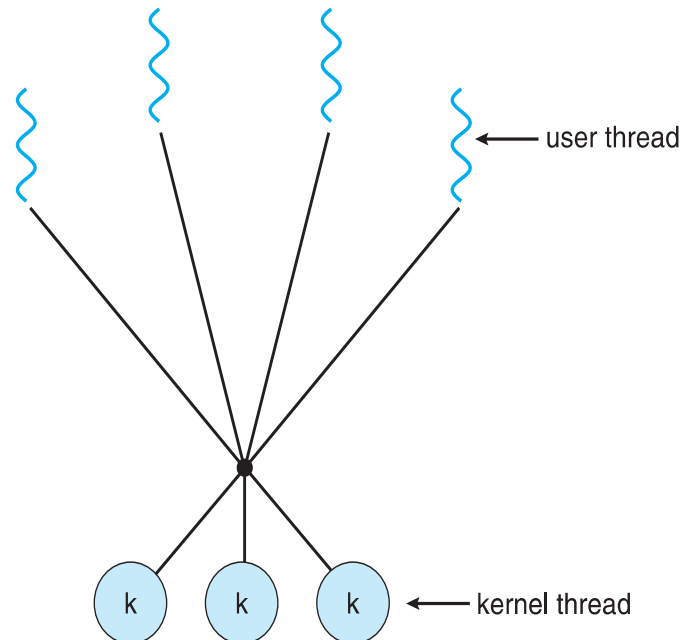


Allows many user level threads to be mapped to many kernel threads

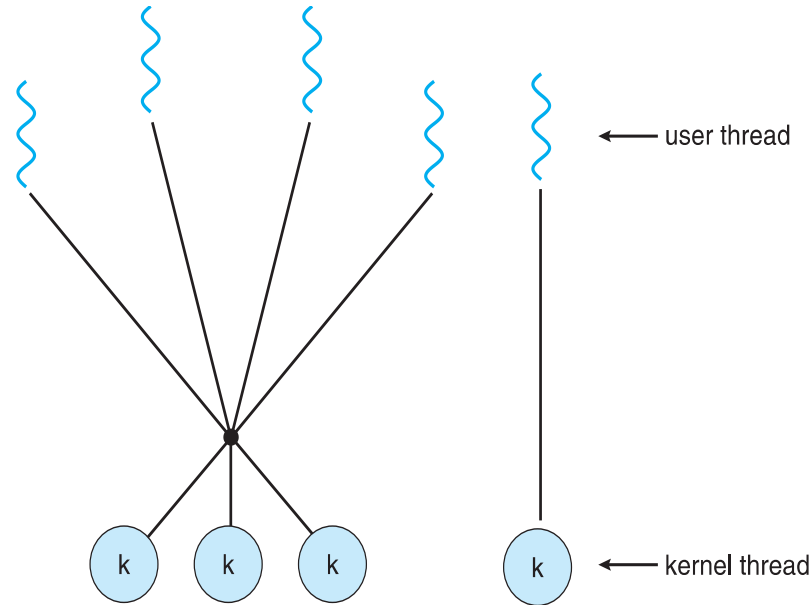
Allows the operating system to create a sufficient number of kernel threads

Solaris prior to version 9

Windows with the *ThreadFiber* package



- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
  - IRIX
  - HP-UX
  - Tru64 UNIX
  - Solaris 8 and earlier





- Distinction between user-level and kernel-level threads
- When threads supported, threads scheduled, not processes
- Many-to-one and many-to-many models, thread library schedules user-level threads to run on LWP
  - Known as **process-contention scope (PCS)** since scheduling competition is within the process
  - Typically done via priority set by programmer
- Kernel thread scheduled onto available CPU is **system-contention scope (SCS)** – competition among all threads in system

- API allows specifying either PCS or SCS during thread creation
  - PTHREAD\_SCOPE\_PROCESS schedules threads using PCS scheduling
  - PTHREAD\_SCOPE\_SYSTEM schedules threads using SCS scheduling
- Can be limited by OS – Linux and Mac OS X only allow PTHREAD\_SCOPE\_SYSTEM

# OPERATING SYSTEMS

## Pthreads Example



```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }

    /* get the default attributes */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

# OPERATING SYSTEMS

## Pthread Scheduling API



```
#include <pthread.h>

#include <stdio.h>

#define NUM_THREADS 5

int main(int argc, char *argv[]) {
    int i, scope;
    pthread_t tid[NUM_THREADS];

    pthread_attr_t attr;

    /* get the default attributes */
    pthread_attr_init(&attr);

    /* first inquire on the current scope */
    if (pthread_attr_getscope(&attr, &scope) != 0)
        fprintf(stderr, "Unable to get scheduling
scope\n");
    else {
        if (scope == PTHREAD_SCOPE_PROCESS)
            printf("PTHREAD_SCOPE_PROCESS");
        else if (scope == PTHREAD_SCOPE_SYSTEM)
            printf("PTHREAD_SCOPE_SYSTEM");
        else
            fprintf(stderr, "Illegal scope value.\n");
    }
    Slides Adapted from Operating System Concepts 9/e © Authors
```

```
/* set the scheduling algorithm
to PCS or SCS */

    pthread_attr_setscope(&attr,
PTHREAD_SCOPE_SYSTEM);

    /* create the threads */
    for (i = 0; i < NUM_THREADS;
i++)

        pthread_create(&tid[i], &attr, run
ner, NULL);

    /* now join on each thread */
    for (i = 0; i < NUM_THREADS;
i++)

        pthread_join(tid[i],
NULL);
}

/* Each thread will begin
control in this function */

void *runner(void *param)
{
    /* do some work ... */
    pthread_exit(0);
}
```

# OPERATING SYSTEMS

## Pthread Scheduling API



```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5

int main(int argc, char *argv[]) {
    int i, scope;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;

    /* get the default attributes */
    pthread_attr_init(&attr);

    /* first inquire on the current scope */
    if (pthread_attr_getscope(&attr, &scope) != 0)
        fprintf(stderr, "Unable to get scheduling
scope\n");
    else {
        if (scope == PTHREAD_SCOPE_PROCESS)
            printf("PTHREAD_SCOPE_PROCESS");
        else if (scope == PTHREAD_SCOPE_SYSTEM)
            printf("PTHREAD_SCOPE_SYSTEM");
        else
            fprintf(stderr, "Illegal scope value.\n");
    }
    Slides Adapted from Operating System Concepts 9/e © Authors
```

```
/* set the scheduling algorithm
to PCS or SCS */

    pthread_attr_setscope(&attr,
PTHREAD_SCOPE_SYSTEM);

    /* create the threads */
    for (i = 0; i < NUM_THREADS;
i++)

pthread_create(&tid[i], &attr, run
ner, NULL);

    /* now join on each thread */
    for (i = 0; i < NUM_THREADS;
i++)

        pthread_join(tid[i],
NULL);
}

/* Each thread will begin
control in this function */

void *runner(void *param)
{
    /* do some work ... */
    pthread_exit(0);
}
```

- An application consists of one or more processes.
- One or more threads run in the context of the process.
- A *job object* allows groups of processes to be managed as a unit.
- Job objects are namable, securable, sharable objects that control attributes of the processes associated with them.
- Operations performed on the job object affect all processes associated with the job object.
- A *thread pool* is a collection of worker threads that efficiently execute asynchronous callbacks on behalf of the application.
- The thread pool is primarily used to reduce the number of application threads and provide management of the worker threads.
- A *fiber* is a unit of execution that must be manually scheduled by the application. Fibers run in the context of the threads that schedule them.
- *User-mode scheduling* (UMS) is a lightweight mechanism that applications can use to schedule their own threads.
- UMS threads differ from fibers in that each UMS thread has its own thread context instead of sharing the thread context of a single thread.

# OPERATING SYSTEMS

## Windows Threading Example



```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }

    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    if (ThreadHandle != NULL) {
        /* now wait for the thread to finish */
        WaitForSingleObject(ThreadHandle, INFINITE);

        /* close the thread handle */
        CloseHandle(ThreadHandle);

        printf("sum = %d\n", Sum);
    }
}
```



**THANK YOU**

---

**Dr Rahul Nagpal**

Computer Science

**[rahulnagpal@pes.edu](mailto:rahulnagpal@pes.edu)**